

## Report 7 — October 29

Lecturer: Wu Qinghua

Completed by: 2022K8009929010 Zhang Jiawei

## 7.1 实验内容

- 在主机上安装 arptables, iptables, 用于禁止每个节点的相应功能。
  - 运行给定网络拓扑: 路由器节点 r1 上执行脚本 (disable\_arp.sh, disable\_icmp.sh, disable\_ip\_forward.sh), 禁止协议栈的相应功能; 终端节点 h1-h3 上执行脚本 disable\_offloading.sh。
  - 在 r1 上执行路由器程序。
  - 在 h1 上进行 ping 实验: Ping 10.0.1.1 (r1), 能够 ping 通; Ping 10.0.2.22 (h2), 能够 ping 通; Ping 10.0.3.33 (h3), 能够 ping 通; Ping 10.0.3.11, 返回 ICMP Destination Host Unreachable; Ping 10.0.4.1, 返回 ICMP Destination Net Unreachable。
- 构造一个包含多个路由器节点组成的网络: 手动配置每个路由器节点的路由表; 有两个终端节点, 通过路由器节点相连, 两节点之间的跳数不少于 3 跳, 手动配置其默认路由表。
  - 连通性测试: 终端节点 ping 每个路由器节点的入端口 IP 地址, 能够 ping 通。
  - 路径测试: 在一个终端节点上 traceroute 另一节点, 能够正确输出路径上每个节点的 IP 信息。

## 7.2 实验过程

### 7.2.1 IP 数据包转发

对于路由器接收到的 IP 报文, 在转发之前需要先进行一次判定。倘若该报文为 ICMP 协议下的请求回复报文 (ICMP\_ECHOREQUEST) 且目的地址为路由器的入端口 IP 地址, 直接向源地址发送 ICMP 回复报文 (ICMP\_ECHOREPLY):

```
struct iphdr *iphdr = packet_to_ip_hdr(packet);
struct icmphdr *icmphdr = (struct icmphdr *)IP_DATA(iphdr);
u32 dest = ntohl(iphdr->daddr);
u8 protocol = iphdr->protocol;
u8 type = icmphdr->type;

// check if the packet is ICMP echo request and the destination IP address is
// equal to the IP address of the iface
if (dest == iface->ip && protocol == IPPROTO_ICMP && type == ICMP_ECHOREQUEST) {
    icmp_send_packet(packet, len, ICMP_ECHOREPLY, 0);
    free(packet);
}
```

```
    return;  
}
```

对于其他情况, 路由器需要根据路由表进行转发。首先, 对 IP 报头的 TTL 字段进行减一操作, 若 TTL 字段减为 0, 则将该数据包丢弃, 并回复 ICMP 信息 (ICMP\_EXC\_TTL)。然后更新校验和, 根据目的 IP 地址查找路由表, 若找到对应的路由表项, 则根据下一跳 IP 地址进行转发; 若未找到对应的路由表项, 则回复 ICMP 信息 (ICMP\_NET\_UNREACH)。下一跳 IP 地址的确定则需要根据路由表项的下一跳网关 (gateway) 地址进行判断, 若网关地址为 0, 说明目的主机在本地网络中, 直接向目的主机发送数据包; 若网关地址不为 0, 说明目的主机在其他网络中, 将数据包发送给网关:

```
// forward the packet  
iphdr->ttl--;  
  
// check if the TTL is less than or equal to 0  
if (iphdr->ttl <= 0) {  
    icmp_send_packet(packet, len, ICMP_TIME_EXCEEDED, ICMP_EXC_TTL);  
    free(packet);  
    return;  
}  
  
// update the checksum of the IP header  
iphdr->checksum = ip_checksum(iphdr);  
  
// search the routing table for the longest prefix match  
rt_entry_t *match = longest_prefix_match(dest);  
if (match == NULL) {  
    icmp_send_packet(packet, len, ICMP_DEST_UNREACH, ICMP_NET_UNREACH);  
    free(packet);  
    return;  
}  
  
// check if the destination IP address is in the same network with the iface  
u32 nextip;  
if (match->gw == 0)  
    nextip = dest;  
else  
    nextip = match->gw;  
iface_send_packet_by_arp(match->iface, nextip, packet, len);
```

还需要说明的是, 查询路由表使用的是最长前缀匹配方法, 若目的地址与路由表中的目的地址匹配, 则比较掩码长度, 选择掩码长度最长的路由表项作为匹配项:

```
// lookup in the routing table, to find the entry with the same and longest  
// prefix.  
// the input address is in host byte order
```

```
rt_entry_t *longest_prefix_match(u32 dst)
{
    // fprintf(stderr, "TODO: longest prefix match for the packet.\n");
    rt_entry_t *entry = NULL, *match = NULL;
    list_for_each_entry(entry, &rtable, list) {
        if ((dst & entry->mask) == (entry->dest & entry->mask)) {
            if (match == NULL || entry->mask > match->mask)
                match = entry;
        }
    }
    return match;
}
```

---

### 7.2.2 ARP 数据包处理

对于路由器接收到的 ARP 报文, 首先分析报文的类型。若为 ARP 请求报文 (ARPOP\_REQUEST), 则判断目的 IP 地址是否为路由器的入端口 IP 地址, 若是则回复 ARP 应答报文, 且将该 ARP 请求报文的源 MAC 地址和 IP 地址添加到 ARP 缓存中; 若为 ARP 回复报文 (ARPOP\_REPLY), 则说明该报文是对路由器发送的 ARP 请求报文的回复, 当回复的目的 IP 地址为路由器的入端口 IP 地址时, 将该 ARP 回复报文的源 MAC 地址和 IP 地址添加到 ARP 缓存中:

```
void handle_arp_packet(iface_info_t *iface, char *packet, int len)
{
    // fprintf(stderr, "TODO: process arp packet: arp request & arp reply.\n");
    struct ether_arp *arp = (struct ether_arp *) (packet + ETHER_HDR_SIZE);

    // arp request
    if (ntohs(arp->arp_op) == ARPOP_REQUEST) {
        if (ntohl(arp->arp_tpa) == iface->ip) {
            arpcache_insert(ntohl(arp->arp_spa), arp->arp_sha);
            arp_send_reply(iface, arp);
        }
    }

    // arp reply
    else if (ntohs(arp->arp_op) == ARPOP_REPLY)
        if (ntohl(arp->arp_tpa) == iface->ip)
            arpcache_insert(ntohl(arp->arp_spa), arp->arp_sha);
    else
        fprintf(stderr, "Unknown ARP packet\n");
    free(packet);
}
```

---

当然,我们也需要发送 ARP 报文,仍然分为 ARP 请求报文和 ARP 回复报文。这两个报文的构造比较类似,首先申请空间然后填入相应的字段,需要注意的是,ARP 请求报文的的目的 MAC 地址为广播地址(全 1),而 ARP 回复报文的的目的 MAC 地址为 ARP 请求报文的源 MAC 地址:

```
// send an arp request: encapsulate an arp request packet, send it out through
// iface_send_packet
void arp_send_request(iface_info_t *iface, u32 dst_ip)
{
    // fprintf(stderr, "TODO: send arp request when lookup failed in arpcache.\n");
    char *packet = (char *)malloc(ETHER_HDR_SIZE + sizeof(struct ether_arp));
    memset(packet, 0, ETHER_HDR_SIZE + sizeof(struct ether_arp));

    struct ether_header *eh = (struct ether_header *)packet;
    memcpy(eh->ether_dhost, "\xff\xff\xff\xff\xff\xff", ETH_ALEN);
    memcpy(eh->ether_shost, iface->mac, ETH_ALEN);
    eh->ether_type = htons(ETH_P_ARP);

    struct ether_arp *arp = (struct ether_arp *)(packet + ETHER_HDR_SIZE);
    arp->arp_hrd = htons(ARPHRD_ETHER);
    arp->arp_pro = htons(ETH_P_IP);
    arp->arp_hln = ETH_ALEN;
    arp->arp_pln = 4;
    arp->arp_op = htons(ARPOP_REQUEST);
    memcpy(arp->arp_sha, iface->mac, ETH_ALEN);
    arp->arp_spa = htonl(iface->ip);
    memset(arp->arp_tha, 0, ETH_ALEN);
    arp->arp_tpa = htonl(dst_ip);

    iface_send_packet(iface, packet, ETHER_HDR_SIZE + sizeof(struct ether_arp));
}

// send an arp reply packet: encapsulate an arp reply packet, send it out
// through iface_send_packet
void arp_send_reply(iface_info_t *iface, struct ether_arp *req_hdr)
{
    // fprintf(stderr, "TODO: send arp reply when receiving arp request.\n");
    char *packet = (char *)malloc(ETHER_HDR_SIZE + sizeof(struct ether_arp));
    struct ether_header *eh = (struct ether_header *)packet;
    memcpy(eh->ether_dhost, req_hdr->arp_sha, ETH_ALEN);
    memcpy(eh->ether_shost, iface->mac, ETH_ALEN);
    eh->ether_type = htons(ETH_P_ARP);

    struct ether_arp *arp = (struct ether_arp *)(packet + ETHER_HDR_SIZE);
```

```
arp->arp_hrd = htons(ARPHRD_ETHER);
arp->arp_pro = htons(ETH_P_IP);
arp->arp_hln = ETH_ALEN;
arp->arp_pln = 4;
arp->arp_op = htons(ARPOP_REPLY);
memcpy(arp->arp_sha, iface->mac, ETH_ALEN);
arp->arp_spa = htonl(iface->ip);
memcpy(arp->arp_tha, req_hdr->arp_sha, ETH_ALEN);
arp->arp_tpa = req_hdr->arp_spa;

iface_send_packet(iface, packet, ETHER_HDR_SIZE + sizeof(struct ether_arp));
}
```

### 7.2.3 ARP 缓存管理

在之前提到的 IP 数据包转发中,最后是通过查找 ARP 缓存来确定下一跳的 MAC 地址。在本实验中,ARP 缓存的管理主要包括查找 ARP 缓存、添加待应答数据包、插入 ARP 缓存和清理 ARP 缓存三个操作。

查找操作是根据目的 IP 地址查找有效 ARP 缓存中是否存在对应的 MAC 地址,若存在则返回 1,否则返回 0:

```
// lookup the IP->mac mapping
//
// traverse the table to find whether there is an entry with the same IP
// and mac address with the given arguments
int arpcache_lookup(u32 ip4, u8 mac[ETH_ALEN])
{
    // fprintf(stderr, "TODO: lookup ip address in arp cache.\n");
    pthread_mutex_lock(&arpcache.lock);
    for (int i = 0; i < MAX_ARP_SIZE; i++) {
        if (arpcache.entries[i].ip4 == ip4 && arpcache.entries[i].valid) {
            memcpy(mac, arpcache.entries[i].mac, ETH_ALEN);
            pthread_mutex_unlock(&arpcache.lock);
            return 1;
        }
    }
    pthread_mutex_unlock(&arpcache.lock);
    return 0;
}
```

但是,ARP 缓存应答并非是立即返回的,而是需要等待 ARP 应答报文的到来。因此,需要添加待应答数据包,即将数据包添加到 ARP 缓存的待应答数据包队列中。这里又分为两种情况,一种是同一 IP 地址的数据包已经在待应答数据包队列中,这说明 ARP 请求报文已经发送过了,只需要将数据包添加到该 IP 地址的待应答

数据包队尾;另一种是同一 IP 地址的数据包不在待应答数据包队列中,这说明 ARP 请求报文还未发送,则需要新建与该 IP 地址对应的待应答数据包队列,然后发送 ARP 请求报文:

```
// append the packet to arpcache
//
// Lookup in the list which stores pending packets, if there is already an
// entry with the same IP address and iface (which means the corresponding arp
// request has been sent out), just append this packet at the tail of that entry
// (the entry may contain more than one packet); otherwise, malloc a new entry
// with the given IP address and iface, append the packet, and send arp request.
void arpcache_append_packet(iface_info_t *iface, u32 ip4, char *packet, int len)
{
    // fprintf(stderr, "TODO: append the ip address if lookup failed, and send arp
    // request if necessary.\n");
    pthread_mutex_lock(&arpcache.lock);
    struct arp_req *req_entry = NULL;
    struct arp_req *req_q;

    // find the entry with the same IP address
    list_for_each_entry_safe(req_entry, req_q, &(arpcache.req_list), list) {
        if (req_entry->ip4 == ip4) {
            struct cached_pkt *pkt_entry = (struct cached_pkt *)malloc(sizeof(struct
                cached_pkt));
            init_list_head(&(pkt_entry->list));
            pkt_entry->packet = packet;
            pkt_entry->len = len;
            list_add_tail(&(pkt_entry->list), &(req_entry->cached_packets));
            pthread_mutex_unlock(&arpcache.lock);
            return;
        }
    }

    // no entry found, malloc a new entry
    req_entry = (struct arp_req *)malloc(sizeof(struct arp_req));
    init_list_head(&(req_entry->list));
    init_list_head(&(req_entry->cached_packets));
    req_entry->iface = iface;
    req_entry->ip4 = ip4;
    req_entry->sent = time(NULL);
    req_entry->retries = 0;
    list_add_tail(&(req_entry->list), &(arpcache.req_list));

    struct cached_pkt *pkt_entry = (struct cached_pkt *)malloc(sizeof(struct
        cached_pkt));
```

```
    pkt_entry->packet = packet;
    pkt_entry->len = len;
    list_add_tail(&(pkt_entry->list), &(req_entry->cached_packets));

    pthread_mutex_unlock(&arpcache.lock);
    arp_send_request(iface, ip4);
}
```

并不是所有的 IP-MAC 映射都存在于缓存中的。当需要向 ARP 缓存中插入新的 ARP 缓存项时,首先查找是否存在相同的 IP 地址,若存在则更新该 ARP 缓存项的 MAC 地址和时间戳;若不存在则查找是否有无效的 ARP 缓存项,若有则直接覆盖内容,倘若所有 ARP 缓存项都是有效的,则取随机条目进行覆盖。之后再查找待应答数据包队列中是否有该 IP 地址的数据包,若有则直接填写目的 MAC 地址并发送:

```
// insert the IP->mac mapping into arpcache, if there are pending packets
// waiting for this mapping, fill the ethernet header for each of them, and send
// them out
void arpcache_insert(u32 ip4, u8 mac[ETH_ALEN])
{
    // fprintf(stderr, "TODO: insert ip->mac entry, and send all the pending
    // packets.\n");
    pthread_mutex_lock(&arpcache.lock);
    int i;

    // update the entry if it already exists
    for (i = 0; i < MAX_ARP_SIZE; i++) {
        if (arpcache.entries[i].valid && arpcache.entries[i].ip4 == ip4){
            arpcache.entries[i].added = time(NULL);
            memcpy(arpcache.entries[i].mac, mac, ETH_ALEN);
            pthread_mutex_unlock(&arpcache.lock);
            return;
        }
    }

    // search for an empty entry
    for (i = 0; i < MAX_ARP_SIZE; i++)
        if (!arpcache.entries[i].valid)
            break;

    // if no empty entry, replace a random entry
    if (i == MAX_ARP_SIZE)
        i = rand() % MAX_ARP_SIZE;

    arpcache.entries[i].ip4 = ip4;
```

```
    arpcache.entries[i].added = time(NULL);
    arpcache.entries[i].valid = 1;
    memcpy(arpcache.entries[i].mac, mac, ETH_ALEN);

    // send all the pending packets with the same IP address of the new entry
    struct arp_req *req_entry = NULL, *req_q;
    list_for_each_entry_safe(req_entry, req_q, &(arpcache.req_list), list) {
        if (req_entry->ip4 == ip4) {
            struct cached_pkt *pkt_entry = NULL, *pkt_q;
            list_for_each_entry_safe(pkt_entry, pkt_q, &(req_entry->cached_packets),
                                    list) {
                memcpy(pkt_entry->packet, mac, ETH_ALEN);
                iface_send_packet(req_entry->iface, pkt_entry->packet,
                                pkt_entry->len);
                list_delete_entry(&(pkt_entry->list));
                free(pkt_entry);
            }
            list_delete_entry(&(req_entry->list));
            free(req_entry);
        }
    }
    pthread_mutex_unlock(&arpcache.lock);
}
```

ARP 缓存也有相应的老化和清理机制,当 ARP 缓存项的时间戳超过 15 秒时,清除该 ARP 缓存项。而且,当每一秒发送一次的 ARP 请求报文的重传次数超过 5 次但仍未收到 ARP 回复报文时,对相应等待队列中的数据包回复 ICMP 无法到达信息 (ICMP\_HOST\_UNREACH) 并删除数据包:

```
// sweep arpcache periodically
//
// For the IP->mac entry, if the entry has been in the table for more than 15
// seconds, remove it from the table.
// For the pending packets, if the arp request is sent out 1 second ago, while
// the reply has not been received, retransmit the arp request. If the arp
// request has been sent 5 times without receiving arp reply, for each
// pending packet, send icmp packet (DEST_HOST_UNREACHABLE), and drop these
// packets.
void *arpcache_sweep(void *arg)
{
    while (1) {
        sleep(1);
        // fprintf(stderr, "TODO: sweep arpcache periodically: remove old entries,
        //         resend arp requests .\n");
    }
}
```



```
pthread_mutex_lock(&arpcache.lock);
for (int i = 0; i < MAX_ARP_SIZE; i++)
    if (arpcache.entries[i].valid && time(NULL) - arpcache.entries[i].added
        > ARP_ENTRY_TIMEOUT)
        arpcache.entries[i].valid = 0;

struct list_head temp;
init_list_head(&temp);

struct arp_req *req_entry = NULL, *req_q;
list_for_each_entry_safe(req_entry, req_q, &(arpcache.req_list), list) {
    if (time(NULL) - req_entry->sent >= 1) {
        req_entry->retries++;
        req_entry->sent = time(NULL);
        if (req_entry->retries > ARP_REQUEST_MAX_RETRIES) {
            struct cached_pkt *pkt_entry = NULL, *pkt_q;
            list_for_each_entry_safe(pkt_entry, pkt_q,
                &(req_entry->cached_packets), list) {
                list_delete_entry(&(pkt_entry->list));
                list_add_tail(&(pkt_entry->list), &temp);
            }
            list_delete_entry(&(req_entry->list));
            free(req_entry);
        }
        else
            arp_send_request(req_entry->iface, req_entry->ip4);
    }
}
pthread_mutex_unlock(&arpcache.lock);

struct cached_pkt *pkt_entry = NULL, *pkt_q;
list_for_each_entry_safe(pkt_entry, pkt_q, &temp, list) {
    icmp_send_packet(pkt_entry->packet, pkt_entry->len, ICMP_DEST_UNREACH,
        ICMP_HOST_UNREACH);
    free(pkt_entry);
}

return NULL;
}
```

## 7.3 ICMP 数据包处理

ICMP 数据包的构造需要花一些功夫。申请存储空间时需要注意数据包的长度,对于非 ICMP\_ECHOREPLY 报文,长度为以太网头 +IP 头 +ICMP 头 + 原数据包;对于 ICMP\_ECHOREPLY 报文,长度计算简化为原数据包长度-原数据包 IP 头长度 +IP 头长度。然后填充相应的字段,对于 ICMP\_ECHOREPLY 报文,源 IP 地址就是数据包的源 IP 地址,其余情况下则查询路由表,找到下一跳网关的 IP 地址并填充;目的 IP 地址则是数据包的源 IP 地址;ICMP 类型和代码直接填充;校验和需要计算。最后即可发送:

```
// send icmp packet
void icmp_send_packet(const char *in_pkt, int len, u8 type, u8 code)
{
    // fprintf(stderr, "TODO: malloc and send icmp packet.\n");
    struct iphdr *iphdr = packet_to_ip_hdr(in_pkt);
    char *ipdata = IP_DATA(iphdr);

    // calculate icmp packet length
    int icmp_len;
    if (type == ICMP_ECHOREPLY)
        icmp_len = ntohs(iphdr->tot_len) - IP_HDR_SIZE(iphdr);
    else
        icmp_len = ICMP_HDR_SIZE + IP_HDR_SIZE(iphdr) + ICMP_COPIED_DATA_LEN;
    int res_len = IP_BASE_HDR_SIZE + ETHER_HDR_SIZE + icmp_len;

    char *res = (char *)malloc(res_len);
    memset(res, 0, res_len);

    // fill ip header
    struct iphdr *res_iphdr = packet_to_ip_hdr(res);
    if (type == ICMP_ECHOREPLY)
        ip_init_hdr(res_iphdr, ntohl(iphdr->daddr), ntohl(iphdr->saddr), icmp_len +
            IP_BASE_HDR_SIZE, IPPROTO_ICMP);
    else{
        rt_entry_t *match = longest_prefix_match(ntohl(iphdr->saddr));
        if (match == NULL) {
            free(res);
            return;
        }
        ip_init_hdr(res_iphdr, match->iface->ip, ntohl(iphdr->saddr), icmp_len +
            IP_BASE_HDR_SIZE, IPPROTO_ICMP);
    }

    char *res_ipdata = IP_DATA(res_iphdr);
```

```
// fill icmp header and data
struct icmphdr *res_icmphdr = (struct icmphdr *)res_ipdata;
if (type == ICMP_ECHOREPLY)
    memcpy(res_ipdata, ipdata, icmp_len);
else
    memcpy(res_ipdata + ICMP_HDR_SIZE, iphdr, icmp_len - ICMP_HDR_SIZE);

res_icmphdr->type = type;
res_icmphdr->code = code;
res_icmphdr->checksum = icmp_checksum(res_icmphdr, icmp_len);

// send icmp packet
ip_send_packet(res, res_len);
}
```

ICMP 数据包的发出过程与发送 IP 数据包类似,都是以同样的方式确定下一跳的 IP 地址,最后以 ARP 机制发送数据包,这里不再赘述。

## 7.4 实验结果

### 7.4.1 给定网络拓扑

实验结果如下:

```
root@zhangjiawei-VirtualBox:/home/zhangjiawei/📁/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.2.22 -c 4
PING 10.0.2.22 (10.0.2.22) 56(84) bytes of data.
64 bytes from 10.0.2.22: icmp_seq=1 ttl=63 time=0.141 ms
64 bytes from 10.0.2.22: icmp_seq=2 ttl=63 time=0.058 ms
64 bytes from 10.0.2.22: icmp_seq=3 ttl=63 time=0.067 ms
64 bytes from 10.0.2.22: icmp_seq=4 ttl=63 time=0.056 ms
```

(a) h1 ping h2

```
root@zhangjiawei-VirtualBox:/home/zhangjiawei/📁/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.3.33 -c 4
PING 10.0.3.33 (10.0.3.33) 56(84) bytes of data.
64 bytes from 10.0.3.33: icmp_seq=1 ttl=63 time=0.106 ms
64 bytes from 10.0.3.33: icmp_seq=2 ttl=63 time=0.059 ms
64 bytes from 10.0.3.33: icmp_seq=3 ttl=63 time=0.057 ms
64 bytes from 10.0.3.33: icmp_seq=4 ttl=63 time=0.067 ms
```

(b) h1 ping h3

可见,h1 能够 ping 通 h2、h3 和 r1,但无法 ping 通其他主机,说明实验结果符合预期。

### 7.4.2 自定义网络拓扑

自定义网络拓扑如下:

```

root@zhangjiawei-VirtualBox:/home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.1.1 -c 4
PING 10.0.1.1 (10.0.1.1) 56(84) bytes of data.
64 bytes from 10.0.1.1: icmp_seq=1 ttl=64 time=0.351 ms
64 bytes from 10.0.1.1: icmp_seq=2 ttl=64 time=0.074 ms
64 bytes from 10.0.1.1: icmp_seq=3 ttl=64 time=0.045 ms
64 bytes from 10.0.1.1: icmp_seq=4 ttl=64 time=0.043 ms

```

(c) h1 ping r1

```

root@zhangjiawei-VirtualBox:/home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.3.11 -c 4
PING 10.0.3.11 (10.0.3.11) 56(84) bytes of data.
From 10.0.1.1 icmp_seq=1 Destination Host Unreachable
From 10.0.1.1 icmp_seq=2 Destination Host Unreachable
From 10.0.1.1 icmp_seq=3 Destination Host Unreachable
From 10.0.1.1 icmp_seq=4 Destination Host Unreachable

```

(d) h1 ping unreachable1

```

root@zhangjiawei-VirtualBox:/home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.4.1 -c 4
PING 10.0.4.1 (10.0.4.1) 56(84) bytes of data.
From 10.0.1.1 icmp_seq=1 Destination Net Unreachable
From 10.0.1.1 icmp_seq=2 Destination Net Unreachable
From 10.0.1.1 icmp_seq=3 Destination Net Unreachable
From 10.0.1.1 icmp_seq=4 Destination Net Unreachable

```

(e) h1 ping unreachable2

图 7.1. 给定网络拓扑实验结果

---

```
h1, h2, r1, r2, r3 = net.get('h1', 'h2', 'r1', 'r2', 'r3')
```

# 配置 IP 地址

h1.cmd('ifconfig h1-eth0 10.0.1.1/24')

h2.cmd('ifconfig h2-eth0 10.0.4.1/24')

r1.cmd('ifconfig r1-eth0 10.0.1.2/24')

r1.cmd('ifconfig r1-eth1 10.0.2.1/24')

r2.cmd('ifconfig r2-eth0 10.0.2.2/24')

r2.cmd('ifconfig r2-eth1 10.0.3.1/24')

r3.cmd('ifconfig r3-eth0 10.0.3.2/24')

r3.cmd('ifconfig r3-eth1 10.0.4.2/24')

# 配置路由表

h1.cmd('route add default gw 10.0.1.2')

h2.cmd('route add default gw 10.0.4.2')

r1.cmd('route add -net 10.0.3.0 netmask 255.255.255.0 gw 10.0.2.2 dev r1-eth1')

r1.cmd('route add -net 10.0.4.0 netmask 255.255.255.0 gw 10.0.2.2 dev r1-eth1')

```
r2.cmd('route add -net 10.0.1.0 netmask 255.255.255.0 gw 10.0.2.1 dev r2-eth0')
r2.cmd('route add -net 10.0.4.0 netmask 255.255.255.0 gw 10.0.3.2 dev r2-eth1')

r3.cmd('route add -net 10.0.1.0 netmask 255.255.255.0 gw 10.0.3.1 dev r3-eth0')
r3.cmd('route add -net 10.0.2.0 netmask 255.255.255.0 gw 10.0.3.1 dev r3-eth0')

# 执行脚本以禁用某些功能
for n in (h1, h2, r1, r2, r3):
    n.cmd('./scripts/disable_offloading.sh')
    n.cmd('./scripts/disable_ipv6.sh')

for r in (r1, r2, r3):
    n.cmd('./scripts/disable_arp.sh')
    n.cmd('./scripts/disable_icmp.sh')
    n.cmd('./scripts/disable_ip_forward.sh')
    n.cmd('./scripts/disable_ipv6.sh')
```

实验结果如下：

```
root@zhangjiawei-VirtualBox:/home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.1.2 -c 4
PING 10.0.1.2 (10.0.1.2) 56(84) bytes of data.
64 bytes from 10.0.1.2: icmp_seq=1 ttl=62 time=1.24 ms
64 bytes from 10.0.1.2: icmp_seq=2 ttl=62 time=0.311 ms
64 bytes from 10.0.1.2: icmp_seq=3 ttl=62 time=0.089 ms
64 bytes from 10.0.1.2: icmp_seq=4 ttl=62 time=0.099 ms
```

(a) h2 ping r1

```
root@zhangjiawei-VirtualBox:/home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
07/07-router# ping 10.0.2.2 -c 4
PING 10.0.2.2 (10.0.2.2) 56(84) bytes of data.
64 bytes from 10.0.2.2: icmp_seq=1 ttl=63 time=0.531 ms
64 bytes from 10.0.2.2: icmp_seq=2 ttl=63 time=0.080 ms
64 bytes from 10.0.2.2: icmp_seq=3 ttl=63 time=0.068 ms
64 bytes from 10.0.2.2: icmp_seq=4 ttl=63 time=0.075 ms
```

(b) h2 ping r2

图 7.2. 自定义网络拓扑 ping 实验结果

```
root@zhangjiawei-VirtualBox:/home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
07/07-router# traceroute 10.0.1.1
traceroute to 10.0.1.1 (10.0.1.1), 30 hops max, 60 byte packets
 1  10.0.4.2 (10.0.4.2)  0.214 ms  0.186 ms  0.181 ms
 2  10.0.3.1 (10.0.3.1)  0.321 ms  0.318 ms  0.314 ms
 3  10.0.2.1 (10.0.2.1)  0.452 ms  0.449 ms  0.446 ms
 4  10.0.1.1 (10.0.1.1)  0.441 ms  0.439 ms  0.434 ms
```

图 7.3. 自定义网络拓扑 traceroute 实验结果

可见, h2 能够 ping 通 r1 和 r2, traceroute 也能够正确输出路径上每个节点的 IP 信息, 说明实验结果符合预期。

## 7.5 实验总结

这是一次任务量较大的实验, 但同时也是一次收获颇丰的实验。通过本次实验, 我学会了如何处理 ARP 数据包、IP 数据包和 ICMP 数据包, 实现了路由器的基本功能。在实验过程中, 我对路由表的查找、ARP 缓存的管理、ICMP 数据包的构造和发送等方面有了更深入的了解, 对网络协议栈的实现有了更深刻的认识。同时, 我也学会了如何构建自定义网络拓扑, 进行 ping 和 traceroute 实验, 对网络通信有了更深入的了解。