

Report 9 — Novmeber 7

Lecturer: Wu Qinghua

Completed by: 2022K8009929010 Zhang Jiawei

9.1 实验内容

1. 实验内容一

- 基于已有代码框架,实现路由器生成和处理 mOSPF Hello/LSU 消息的相关操作,构建一致性链路状态数据库;
- 运行实验
 - 运行网络拓扑 (topo.py)
 - 在各个路由器节点上执行 disable_arp.sh, disable_icmp.sh, disable_ip_forward.sh, 禁止协议栈的相应功能
 - 运行 ./mospfd,使得各个节点生成一致的链路状态数据库。

2. 实验内容二

- 基于实验一,实现路由器计算路由表项的相关操作;
- 运行实验
 - 运行网络拓扑 (topo.py)
 - 在各个路由器节点上执行 disable_arp.sh, disable_icmp.sh, disable_ip_forward.sh, 禁止协议栈的相应功能
 - 运行 ./mospfd,使得各个节点生成一致的链路状态数据库
 - 等待一段时间后,每个节点生成完整的路由表项
 - 在节点 h1 上 ping/traceroute 节点 h2
 - 关掉某节点或链路,等一段时间后,再次用 h1 去 traceroute 节点 h2。

9.2 实验过程

9.2.1 链路邻居发现

对于每一个链路上的路由器节点,都会从各个端口以 5 秒为周期发送 Hello 消息,以便发现邻居节点。Hello 消息包括了发送节点的 ID、端口的子网掩码等信息。需要注意一点,Hello 消息的目的 IP 地址是 224.0.0.5,目的 MAC 地址是 01:00:5e:00:00:05,这是因为 Hello 消息是多播消息,需要发送给多播组。再接上各层头部,即可发送。具体代码如下:

```
void *sending_mospf_hello_thread(void *param)
{
    // fprintf(stdout, "TODO: send mOSPF Hello message periodically.\n");
    while (1)
    {
        sleep(MOSPF_DEFAULT_HELLOINT);
        pthread_mutex_lock(&mospf_lock);

        iface_info_t *iface = NULL;
        list_for_each_entry(iface, &instance->iface_list, list) {
            int mospf_len = MOSPF_HDR_SIZE + MOSPF_HELLO_SIZE;
            int packet_len = ETHER_HDR_SIZE + IP_BASE_HDR_SIZE + mospf_len;
            char *packet = (char *)malloc(packet_len);
            struct ether_header *ethhdr = (struct ether_header *)packet;
            struct iphdr *iphdr = packet_to_ip_hdr(packet);
            struct mospf_hdr *mospfhdr = (struct mospf_hdr *)((char *)iphdr +
                IP_BASE_HDR_SIZE);
            struct mospf_hello *hello = (struct mospf_hello *)((char *)mospfhdr +
                MOSPF_HDR_SIZE);

            memset(packet, 0, packet_len);

            mospf_init_hello(hello, iface->mask);
            mospf_init_hdr(mospfhdr, MOSPF_TYPE_HELLO, mospf_len,
                instance->router_id, instance->area_id);
            mospfhdr->checksum = mospf_checksum(mospfhdr);

            ip_init_hdr(iphdr, iface->ip, MOSPF_ALLSPFRouters, packet_len -
                ETHER_HDR_SIZE, IPPROTO_MOSPF);

            memcpy(ethhdr->ether_dhost, "\x01\x00\x5e\x00\x00\x05", ETH_ALEN);
            memcpy(ethhdr->ether_shost, iface->mac, ETH_ALEN);
            ethhdr->ether_type = htons(ETH_P_IP);

            iface_send_packet(iface, packet, packet_len);
        }

        pthread_mutex_unlock(&mospf_lock);
    }

    return NULL;
}
```

当某个路由器节点接收到邻居节点发来的 Hello 消息时, 会根据 Hello 消息中的信息更新邻居节点的信息。若源节点在自身邻居列表中, 则更新对应存活时间等等数据项; 若源节点不在自身邻居列表中, 则将其加入邻居列表, 导致邻居列表变动, 所以还需要发送 LSU 消息通知其他节点。具体代码如下:

```
void handle_mospf_hello(iface_info_t *iface, const char *packet, int len)
{
    // fprintf(stdout, "TODO: handle mOSPF Hello message.\n");
    struct iphdr *iphdr = (struct iphdr *) (packet + ETHER_HDR_SIZE);
    struct mospf_hdr *mospfhdr = (struct mospf_hdr *) ((char *) iphdr +
        IP_HDR_SIZE(iphdr));
    struct mospf_hello *hello = (struct mospf_hello *) ((char *) mospfhdr +
        MOSPF_HDR_SIZE);
    pthread_mutex_lock(&mospf_lock);

    mospf_nbr_t *nbr = NULL;
    list_for_each_entry(nbr, &iface->nbr_list, list) {
        if (nbr->nbr_id == ntohl(mospfhdr->rid)) {
            nbr->alive = 0;
            nbr->nbr_ip = ntohl(iphdr->saddr);
            nbr->nbr_mask = ntohl(hello->mask);

            pthread_mutex_unlock(&mospf_lock);
            return;
        }
    }

    nbr = (mospf_nbr_t *) malloc(sizeof(mospf_nbr_t));
    nbr->nbr_id = ntohl(mospfhdr->rid);
    nbr->nbr_ip = ntohl(iphdr->saddr);
    nbr->nbr_mask = ntohl(hello->mask);
    nbr->alive = 0;
    init_list_head(&nbr->list);
    list_add_tail(&nbr->list, &iface->nbr_list);
    iface->num_nbr++;

    send_mospf_lsu_packet();
    pthread_mutex_unlock(&mospf_lock);
}
```

此外, 邻居列表还需要一个老化操作, 如果列表中的节点在 $3 \times \text{hello-interval}$ 时间内未更新, 则将其删除, 也导致邻居列表变动, 需要发送 LSU 消息通知其他节点。具体代码如下:

```
void *checking_nbr_thread(void *param)
{
```

```
// fprintf(stdout, "TODO: neighbor list timeout operation.\n");
while (1)
{
    sleep(1);
    pthread_mutex_lock(&mospf_lock);

    int delete = 0;
    iface_info_t *iface = NULL;
    list_for_each_entry(iface, &instance->iface_list, list) {
        mospf_nbr_t *nbr = NULL;
        mospf_nbr_t *nbr_q = NULL;
        list_for_each_entry_safe(nbr, nbr_q, &iface->nbr_list, list) {
            if (nbr->alive > 3 * iface->helloint){
                delete = 1;
                iface->num_nbr--;
                list_delete_entry(&nbr->list);
                free(nbr);
            }
            else
                nbr->alive++;
        }
    }

    if (delete)
        send_mospf_lsu_packet();

    pthread_mutex_unlock(&mospf_lock);
}

return NULL;
}
```

9.2.2 链路状态洪泛

路由器节点每隔一段时间就会发送一次 LSU 消息,以便通知其他节点自身的链路状态。我将发送过程封装成一个函数,发送完成后打印节点数据库信息。具体代码如下:

```
void *sending_mospf_lsu_thread(void *param)
{
    // fprintf(stdout, "TODO: send mOSPF LSU message periodically.\n");
    while (1)
    {
        sleep(MOSPF_DEFAULT_LSUINT);
```

```
    pthread_mutex_lock(&mospf_lock);
    send_mospf_lsu_packet();
    print_database();
    pthread_mutex_unlock(&mospf_lock);
}
return NULL;
}
```

发送函数的具体实现如下,首先生成要发送的 LSA 部分,LSA 消息的个数为所有邻居节点个数加上无邻居节点的接口个数,内容包括网络地址、子网掩码、邻居节点 ID(无邻居则为 0)。然后生成 LSU 消息,对于各个端口,复制其相应的 LSA 消息,添加头部信息,向邻居节点发送。具体代码如下:

```
void send_mospf_lsu_packet()
{
    int lsa_num = 0;
    iface_info_t *iface = NULL;
    list_for_each_entry(iface, &instance->iface_list, list) {
        lsa_num += iface->num_nbr ? iface->num_nbr : 1;
    }

    int mospf_len = MOSPF_HDR_SIZE + MOSPF_LSU_SIZE + lsa_num * MOSPF_LSA_SIZE;
    char *packet = (char *)malloc(mospf_len);
    memset(packet, 0, mospf_len);
    struct mospf_hdr *mospfhdr = (struct mospf_hdr *)packet;
    struct mospf_lsu *lsu = (struct mospf_lsu *)((char *)mospfhdr +
        MOSPF_HDR_SIZE);
    struct mospf_lsa *lsa = (struct mospf_lsa *)((char *)lsu + MOSPF_LSU_SIZE);

    iface = NULL;
    list_for_each_entry(iface, &instance->iface_list, list) {
        if (iface->num_nbr) {
            mospf_nbr_t *nbr = NULL;
            list_for_each_entry(nbr, &iface->nbr_list, list) {
                lsa->network = iface->ip & iface->mask;
                lsa->mask = iface->mask;
                lsa->rid = nbr->nbr_id;
                lsa++;
            }
        }
        else {
            lsa->network = iface->ip & iface->mask;
            lsa->mask = iface->mask;
            lsa->rid = 0;
        }
    }
}
```

```

        lsa++;
    }
}

mospf_init_lsu(lsu, lsa_num);
instance->sequence_num++;

mospf_init_hdr(mospfhdr, MOSPF_TYPE_LSU, mospf_len, instance->router_id,
    instance->area_id);
mospfhdr->checksum = mospf_checksum(mospfhdr);

iface = NULL;
list_for_each_entry(iface, &instance->iface_list, list) {
    if (iface->num_nbr) {
        mospf_nbr_t *nbr = NULL;
        list_for_each_entry(nbr, &iface->nbr_list, list) {
            char *packet_send = (char *)malloc(ETHER_HDR_SIZE + IP_BASE_HDR_SIZE
                + mospf_len);
            struct ether_header *ethhdr = (struct ether_header *)packet_send;
            struct iphdr *iphdr = (struct iphdr *)(packet_send + ETHER_HDR_SIZE);
            char *mospf_message = packet_send + ETHER_HDR_SIZE + IP_BASE_HDR_SIZE;

            memset(packet_send, 0, ETHER_HDR_SIZE + IP_BASE_HDR_SIZE + mospf_len);
            memcpy(mospf_message, packet, mospf_len);
            ip_init_hdr(iphdr, iface->ip, nbr->nbr_ip, mospf_len +
                IP_BASE_HDR_SIZE, IPPROTO_MOSPF);
            memcpy(ethhdr->ether_shost, iface->mac, ETH_ALEN);
            ethhdr->ether_type = htons(ETH_P_IP);

            iface_send_packet_by_arp(iface, nbr->nbr_ip, packet_send,
                ETHER_HDR_SIZE + IP_BASE_HDR_SIZE + mospf_len);
        }
    }
}
free(packet);
}

```

当某个路由器节点接收到 LSU 消息时,先检查是否是自身发送的,若是则丢弃。若不是,则根据 LSA 消息更新链路状态数据库,如果数据库内有相同 rid 的条目且新序列号大于数据库内的序列号,则更新,如果数据库内的序列号较大,则丢弃;若数据库内无相同 rid 的条目,则加入数据库。对于洪泛操作,LSU 消息也会转发到邻居节点,如果链路状态得到更新,则检查 LSU 消息的 ttl,若 ttl 大于 0,则继续转发给邻居节点并更新路由表。具体代码如下:

```
void handle_mospf_lsu(iface_info_t *iface, char *packet, int len)
{
    // fprintf(stdout, "TODO: handle mOSPF LSU message.\n");
    struct ether_header *ethhdr = (struct ether_header *)packet;
    struct iphdr *iphdr = (struct iphdr *)(packet + ETHER_HDR_SIZE);
    struct mospf_hdr *mospfhdr = (struct mospf_hdr *)((char *)iphdr +
        IP_HDR_SIZE(iphdr));
    struct mospf_lsu *lsu = (struct mospf_lsu *)((char *)mospfhdr +
        MOSPF_HDR_SIZE);
    struct mospf_lsa *lsa = (struct mospf_lsa *)((char *)lsu + MOSPF_LSU_SIZE);

    pthread_mutex_lock(&mospf_lock);
    if (instance->router_id == ntohl(mospfhdr->rid)) {
        pthread_mutex_unlock(&mospf_lock);
        return;
    }

    int update = 0;
    int exist = 0;
    mospf_db_entry_t *db = NULL;
    list_for_each_entry(db, &mospf_db, list) {
        if (db->rid == ntohl(mospfhdr->rid)) {
            exist = 1;
            if (db->seq < ntohs(lsu->seq)) {
                db->seq = ntohs(lsu->seq);
                db->alive = 0;
                db->nadv = ntohl(lsu->nadv);
                for (int i=0; i<db->nadv; i++) {
                    db->array[i].network = lsa[i].network;
                    db->array[i].mask = lsa[i].mask;
                    db->array[i].rid = lsa[i].rid;
                }
                update = 1;
            }
        }
    }

    if (!exist) {
        db = (mospf_db_entry_t *)malloc(sizeof(mospf_db_entry_t));
        db->rid = ntohl(mospfhdr->rid);
        db->seq = ntohs(lsu->seq);
        db->alive = 0;
        db->nadv = ntohl(lsu->nadv);
    }
}
```

```
db->array = (struct mospf_lsa *)malloc(MOSPF_LSA_SIZE * db->nadv);
for (int i=0; i<db->nadv; i++) {
    db->array[i].network = lsa[i].network;
    db->array[i].mask = lsa[i].mask;
    db->array[i].rid = lsa[i].rid;
}
init_list_head(&db->list);
list_add_tail(&db->list, &mospf_db);
update = 1;
}

if (!update) {
    pthread_mutex_unlock(&mospf_lock);
    return;
}

lsu->ttl--;
if (lsu->ttl > 0) {
    mospfhdr->checksum = mospf_checksum(mospfhdr);
    iface_info_t *iface_out = NULL;

    list_for_each_entry(iface_out, &instance->iface_list, list) {
        if (iface_out->num_nbr && iface_out != iface) {
            mospf_nbr_t *nbr = NULL;
            list_for_each_entry(nbr, &iface_out->nbr_list, list) {
                char *packet_send = (char *)malloc(len);
                struct ether_header *ethdr_send = (struct ether_header
                    *)packet_send;
                struct iphdr *iphdr_send = (struct iphdr *)(packet_send +
                    ETHER_HDR_SIZE);
                memcpy(packet_send, packet, len);

                iphdr_send->daddr = htonl(nbr->nbr_ip);
                iphdr_send->checksum = ip_checksum(iphdr_send);

                memcpy(ethdr_send->ether_shost, iface_out->mac, ETH_ALEN);
                ethdr_send->ether_type = htons(ETH_P_IP);

                iface_send_packet_by_arp(iface_out, nbr->nbr_ip, packet_send, len);
            }
        }
    }
}
```



```
update_rtable();  
pthread_mutex_unlock(&mospf_lock);  
}
```

同样也需要处理节点失效问题,当数据库中一个节点的链路状态超过 40 秒未更新时,表明该节点已失效,将对应条目删除,再更新路由表。具体代码如下:

```
void *checking_database_thread(void *param)  
{  
    // fprintf(stdout, "TODO: link state database timeout operation.\n");  
    while (1)  
    {  
        sleep(1);  
        pthread_mutex_lock(&mospf_lock);  
  
        int delete = 0;  
        mospf_db_entry_t *db = NULL;  
        mospf_db_entry_t *db_q = NULL;  
        list_for_each_entry_safe(db, db_q, &mospf_db, list) {  
            if (db->alive > MOSPF_DATABASE_TIMEOUT){  
                delete = 1;  
                list_delete_entry(&db->list);  
                free(db->array);  
                free(db);  
            }  
            else  
                db->alive++;  
        }  
  
        if (delete)  
            update_rtable();  
  
        pthread_mutex_unlock(&mospf_lock);  
    }  
  
    return NULL;  
}
```

9.2.3 路由表更新计算

路由表更新计算在本实验中是最麻烦的一个步骤。我通过邻接矩阵记录图的拓扑结构,通过 Dijkstra 算法计算最短路径,再根据最短路径更新路由表。具体来说,先清除路由表,但是保留默认路由表项,然后根据节点的链路状态数据库,初始化邻接矩阵,再使用 Dijkstra 算法计算最短路径,最后根据最短路径更新路由表。具体

代码如下:

```
void update_rtable(void)
{
    // clear the routing table, but keep the default entry whose gw is 0
    rt_entry_t *rt_entry, *rt_q;
    list_for_each_entry_safe(rt_entry, rt_q, &rtable, list) {
        if (rt_entry->gw)
            remove_rt_entry(rt_entry);
    }

    // number all the nodes
    // 0 is the router itself
    // 1 ~ n are the other nodes
    node_map[0] = instance->router_id;
    node_num = 1;
    iface_info_t *iface = NULL;
    list_for_each_entry(iface, &instance->iface_list, list) {
        if (iface->num_nbr) {
            mospf_nbr_t *nbr = NULL;
            list_for_each_entry(nbr, &iface->nbr_list, list) {
                if (!rid_exist(nbr->nbr_id))
                    node_map[node_num++] = nbr->nbr_id;
            }
        }
    }

    mospf_db_entry_t *db = NULL;
    list_for_each_entry(db, &mospf_db, list) {
        if (!rid_exist(db->rid))
            node_map[node_num++] = db->rid;
        for (int i=0; i<db->nadv; i++)
            if (db->array[i].rid && !rid_exist(db->array[i].rid))
                node_map[node_num++] = db->array[i].rid;
    }

    // initialize the graph
    for (int i=0; i<node_num; i++) {
        for (int j=0; j<node_num; j++) {
            if (i == j)
                graph[i][j] = 0;
            else
                graph[i][j] = 0x1fffffff; // seen as the maximum value
        }
    }
}
```

```
}

iface = NULL;
list_for_each_entry(iface, &instance->iface_list, list) {
    if (iface->num_nbr) {
        mospf_nbr_t *nbr = NULL;
        list_for_each_entry(nbr, &iface->nbr_list, list) {
            graph[0][get_node_index(nbr->nbr_id)] = 1;
            graph[get_node_index(nbr->nbr_id)][0] = 1;
        }
    }
}

db = NULL;
list_for_each_entry(db, &mospf_db, list) {
    int src = get_node_index(db->rid);
    for (int i=0; i<db->nadv; i++) {
        if (db->array[i].rid) {
            int dst = get_node_index(db->array[i].rid);
            graph[src][dst] = 1;
            graph[dst][src] = 1;
        }
    }
}

// Dijkstra algorithm
int dist[MAX_NODE_NUM];
int visited[MAX_NODE_NUM];
for (int i=0; i<node_num; i++) {
    dist[i] = 0xffffffff;
    visited[i] = 0;
    prev[i] = -1;
}
dist[0] = 0;
stack_top = 0;

for (int i=0; i<node_num; i++) {
    int min = 0xffffffff;
    int u = -1;
    for (int j=0; j<node_num; j++) {
        if (!visited[j] && dist[j] < min) {
            min = dist[j];
            u = j;
        }
    }
}
```

```
    }  
}  
  
if (u == -1)  
    break;  
  
visited[u] = 1;  
stack[stack_top++] = u;  
  
for (int v=0; v<node_num; v++) {  
    if (!visited[v] && dist[u] + graph[u][v] < dist[v]) {  
        dist[v] = dist[u] + graph[u][v];  
        prev[v] = u;  
    }  
}  
}  
  
// update the routing table  
// use the prev array to find the next hop  
db = NULL;  
iface_info_t *iface_out = NULL;  
int current_node = 0;  
int find = 0;  
u32 gw;  
for (int i = 1; i < stack_top; i++) {  
    current_node = stack[i];  
    mospf_db_entry_t *db_tmp = NULL;  
    list_for_each_entry(db_tmp, &mospf_db, list) {  
        if (db_tmp->rid == node_map[current_node]) {  
            db = db_tmp;  
            break;  
        }  
    }  
}  
  
if (db == NULL)  
    continue;  
  
while (prev[current_node] != 0)  
    current_node = prev[current_node];  
  
iface = NULL;  
list_for_each_entry(iface, &instance->iface_list, list) {  
    if (iface->num_nbr) {
```

```
    mospf_nbr_t *nbr = NULL;
    list_for_each_entry(nbr, &iface->nbr_list, list) {
        if (nbr->nbr_id == node_map[current_node]) {
            iface_out = iface;
            gw = nbr->nbr_ip;
            break;
        }
    }
}

if (iface_out == NULL)
    continue;

for (int j = 0; j < db->nadv; j++) {
    find = 0;
    rt_entry_t *rt_entry = NULL;
    list_for_each_entry(rt_entry, &rtable, list) {
        if (rt_entry->dest == db->array[j].network) {
            find = 1;
            break;
        }
    }

    if (!find) {
        rt_entry = new_rt_entry(db->array[j].network, db->array[j].mask, gw,
            iface_out);
        add_rt_entry(rt_entry);
    }
}

print_rtable();

return;
}
```

9.3 实验结果

在每个路由器节点上运行./mospfd,使得各个节点生成一致的链路状态数据库,并查看四个节点的路由表,如下:

"Node: r1">

10.0.4.4	10.0.4.0	255.255.255.0	10.0.2.2
10.0.4.4	10.0.5.0	255.255.255.0	10.0.3.3
10.0.4.4	10.0.6.0	255.255.255.0	0.0.0.0

Routing Table:

dest	mask	gateway	if_name
10.0.1.0	255.255.255.0	0.0.0.0	r1-eth0
10.0.2.0	255.255.255.0	0.0.0.0	r1-eth1
10.0.3.0	255.255.255.0	0.0.0.0	r1-eth2
10.0.4.0	255.255.255.0	10.0.2.2	r1-eth1
10.0.5.0	255.255.255.0	10.0.3.3	r1-eth2
10.0.6.0	255.255.255.0	10.0.2.2	r1-eth1

Routing Table:

dest	mask	gateway	if_name
10.0.1.0	255.255.255.0	0.0.0.0	r1-eth0
10.0.2.0	255.255.255.0	0.0.0.0	r1-eth1
10.0.3.0	255.255.255.0	0.0.0.0	r1-eth2
10.0.4.0	255.255.255.0	10.0.2.2	r1-eth1
10.0.5.0	255.255.255.0	10.0.3.3	r1-eth2
10.0.6.0	255.255.255.0	10.0.2.2	r1-eth1

图 9.1. r1 的数据库和路由表

"Node: r2">

MOSPF Database:

RID	Network	Mask	Neighbor
10.0.4.4	10.0.4.0	255.255.255.0	10.0.2.2
10.0.4.4	10.0.5.0	255.255.255.0	10.0.3.3
10.0.4.4	10.0.6.0	255.255.255.0	0.0.0.0
10.0.1.1	10.0.1.0	255.255.255.0	0.0.0.0
10.0.1.1	10.0.2.0	255.255.255.0	10.0.2.2
10.0.1.1	10.0.3.0	255.255.255.0	10.0.3.3
10.0.3.3	10.0.3.0	255.255.255.0	10.0.1.1
10.0.3.3	10.0.5.0	255.255.255.0	10.0.4.4

Routing Table:

dest	mask	gateway	if_name
10.0.2.0	255.255.255.0	0.0.0.0	r2-eth0
10.0.4.0	255.255.255.0	0.0.0.0	r2-eth1
10.0.1.0	255.255.255.0	10.0.2.1	r2-eth0
10.0.3.0	255.255.255.0	10.0.2.1	r2-eth0
10.0.5.0	255.255.255.0	10.0.4.4	r2-eth1
10.0.6.0	255.255.255.0	10.0.4.4	r2-eth1

图 9.2. r2 的数据库和路由表

9-14

"Node: r3"			
MOSPF Database:			
RID	Network	Mask	Neighbor
10.0.4.4	10.0.4.0	255.255.255.0	10.0.2.2
10.0.4.4	10.0.5.0	255.255.255.0	10.0.3.3
10.0.4.4	10.0.6.0	255.255.255.0	0.0.0.0
10.0.1.1	10.0.1.0	255.255.255.0	0.0.0.0
10.0.1.1	10.0.2.0	255.255.255.0	10.0.2.2
10.0.1.1	10.0.3.0	255.255.255.0	10.0.3.3
10.0.2.2	10.0.2.0	255.255.255.0	10.0.1.1
10.0.2.2	10.0.4.0	255.255.255.0	10.0.4.4
Routing Table:			
dest	mask	gateway	if_name
10.0.3.0	255.255.255.0	0.0.0.0	r3-eth0
10.0.5.0	255.255.255.0	0.0.0.0	r3-eth1
10.0.1.0	255.255.255.0	10.0.3.1	r3-eth0
10.0.2.0	255.255.255.0	10.0.3.1	r3-eth0
10.0.4.0	255.255.255.0	10.0.5.4	r3-eth1
10.0.6.0	255.255.255.0	10.0.5.4	r3-eth1

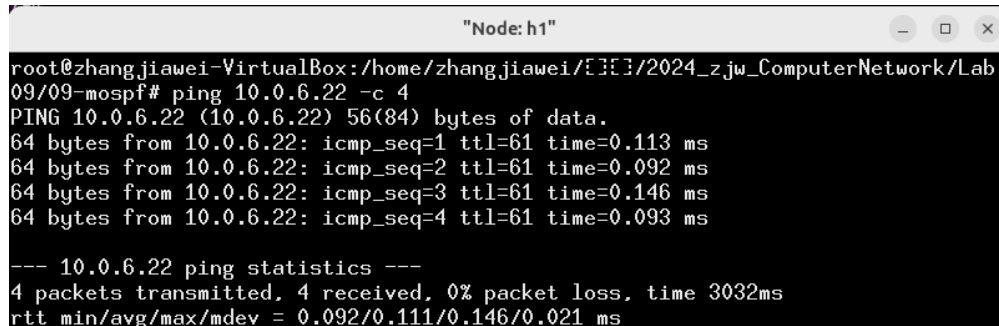
图 9.3. r3 的数据库和路由表

"Node: r4"			
MOSPF Database:			
RID	Network	Mask	Neighbor
10.0.2.2	10.0.2.0	255.255.255.0	10.0.1.1
10.0.2.2	10.0.4.0	255.255.255.0	10.0.4.4
10.0.3.3	10.0.3.0	255.255.255.0	10.0.1.1
10.0.3.3	10.0.5.0	255.255.255.0	10.0.4.4
10.0.1.1	10.0.1.0	255.255.255.0	0.0.0.0
10.0.1.1	10.0.2.0	255.255.255.0	10.0.2.2
10.0.1.1	10.0.3.0	255.255.255.0	10.0.3.3
Routing Table:			
dest	mask	gateway	if_name
10.0.4.0	255.255.255.0	0.0.0.0	r4-eth0
10.0.5.0	255.255.255.0	0.0.0.0	r4-eth1
10.0.6.0	255.255.255.0	0.0.0.0	r4-eth2
10.0.2.0	255.255.255.0	10.0.4.2	r4-eth0
10.0.3.0	255.255.255.0	10.0.5.3	r4-eth1
10.0.1.0	255.255.255.0	10.0.4.2	r4-eth0

图 9.4. r4 的数据库和路由表

可以看出,四个节点的数据库一致,路由表项均存在且正确。

在节点 h1 上 ping 节点 h2,如下:

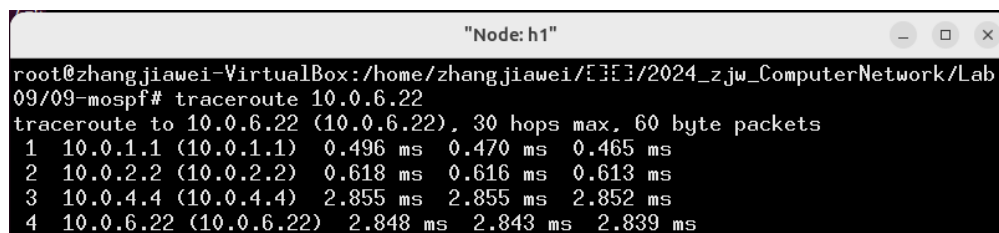


```
root@zhangjiawei-VirtualBox: /home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
09/09-mospf# ping 10.0.6.22 -c 4
PING 10.0.6.22 (10.0.6.22) 56(84) bytes of data:
64 bytes from 10.0.6.22: icmp_seq=1 ttl=61 time=0.113 ms
64 bytes from 10.0.6.22: icmp_seq=2 ttl=61 time=0.092 ms
64 bytes from 10.0.6.22: icmp_seq=3 ttl=61 time=0.146 ms
64 bytes from 10.0.6.22: icmp_seq=4 ttl=61 time=0.093 ms

--- 10.0.6.22 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3032ms
rtt min/avg/max/mdev = 0.092/0.111/0.146/0.021 ms
```

图 9.5. h1 ping h2

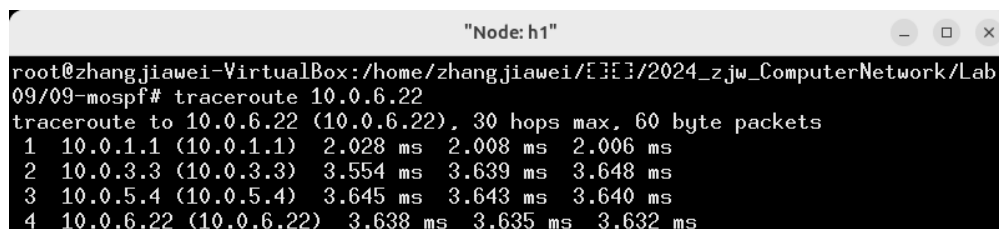
在节点 h1 上 traceroute 节点 h2,如下:



```
root@zhangjiawei-VirtualBox: /home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
09/09-mospf# traceroute 10.0.6.22
traceroute to 10.0.6.22 (10.0.6.22), 30 hops max, 60 byte packets
 1 10.0.1.1 (10.0.1.1) 0.496 ms 0.470 ms 0.465 ms
 2 10.0.2.2 (10.0.2.2) 0.618 ms 0.616 ms 0.613 ms
 3 10.0.4.4 (10.0.4.4) 2.855 ms 2.855 ms 2.852 ms
 4 10.0.6.22 (10.0.6.22) 2.848 ms 2.843 ms 2.839 ms
```

图 9.6. h1 traceroute h2

关闭 r1 和 r2 之间的链路,等待一段时间后,再次用 h1 去 traceroute 节点 h2,如下:



```
root@zhangjiawei-VirtualBox: /home/zhangjiawei/2024_zjw_ComputerNetwork/Lab
09/09-mospf# traceroute 10.0.6.22
traceroute to 10.0.6.22 (10.0.6.22), 30 hops max, 60 byte packets
 1 10.0.1.1 (10.0.1.1) 2.028 ms 2.008 ms 2.006 ms
 2 10.0.3.3 (10.0.3.3) 3.554 ms 3.639 ms 3.648 ms
 3 10.0.5.4 (10.0.5.4) 3.645 ms 3.643 ms 3.640 ms
 4 10.0.6.22 (10.0.6.22) 3.638 ms 3.635 ms 3.632 ms
```

图 9.7. h1 traceroute h2 after link down

路径发生了变化,说明路由表更新成功,算法正确。

9.4 实验总结

本次实验主要是实现 mOSPF 协议的 Hello/LSU 消息的生成和处理,以及路由表的更新计算。实验中需要注意的是,链路状态数据库的一致性,需要在每个节点上都运行 mospfd,以便生成一致的数据库;路由表的更新计算,需要使用 Dijkstra 算法计算最短路径,再根据最短路径更新路由表。实验中还需要注意处理节点失效

问题,当链路状态超过 40 秒未更新时,表明该节点已失效,需要删除对应条目,再更新路由表。我学会了如何实现 mOSPF 协议的 Hello/LSU 消息的生成和处理,以及路由表的更新计算,对于理解路由器的工作原理有很大帮助。